

Stat 405/705
Class 5
Statistical computing with R

Richard P. Waterman

Wharton

Table of contents I

- 1 Today's module
- 2 Last time
- 3 Writing functions
- 4 Summary
- 5 Next time

Today's module

Topics to be covered in this module:

- Last time
- Prediction in `lm` and `glm`
- Writing functions
- Functions used in today's class
- Next time

- Linking R to a database
- Analytics engine architecture
- Joining/merging data frames
- Linear and logistic regression in R: the `lm` command

Generic functions

In both the `lm` and `glm` contexts we can use the same set of functions:

- `summary`
- `predict`
- `residuals`
- `plot`

These are so called generic functions. They look to see the class of the object that they are operating on, and then call the appropriate specific version of the generic function.

Writing functions

- Functions let you organize your code into re-usable chunks
- Functions have three parts:
 - ① The arguments to the function (`formals`)
 - ② The R-code that implements the function itself (`body`)
 - ③ The environment in which the function's variables are found (`environment`)
- The R command to create a function is `function`. Sometimes we write the functions ourselves, often we find someone else's implementation and make use of that (there are probably a couple million functions available in R by now).

Writing functions

The R syntax to write a function:

```
# A function to calculate the percent difference between two numbers
#
percdiff <- function(first,second){
  result <- (second - first)/first * 100 # The calculation itself

#The "sprintf" function gives you complete control over output formatting

#Three places of decimals followed by a % sign
#It will return a "character" so not useful for further calculations!

final.result <- sprintf("%.3f%%",result)
return(final.result) #Explicitly return the result
}
percdiff(first = 5, second = 4) # Now call the function

## [1] "-20.000%"
```

What happened?

- The function named `percdiff` has two "slots" to fill.
- The names of these slots are `first` and `second`.
- The function is written in terms of the names of the slots.
- When the function is evaluated: `"percdiff(first = 5, second = 4)"` the slots are replaced with whatever is in the argument values (5 and 4 in this case) and the result returned from the function.
- Note: If you don't put a `return` statement in the function, then the function will send back the result of the last line of evaluation in the body of the function.
- Suggestion: always include a `return` statement explicitly.

The pieces of this function

```
body(percdiff) # The body of the function
```

```
## {  
##     result <- (second - first)/first * 100  
##     final.result <- sprintf("%.3f%", result)  
##     return(final.result)  
## }
```

```
formals(percdiff) # The arguments to the function
```

```
## $first  
##  
##  
## $second
```

```
environment(percdiff) # The environment in which the function was called
```

```
## <environment: R_GlobalEnv>
```

Calling functions

#There are many ways to call this function:

```
percdiff(first = 3, second = 4) # Naming the arguments explicitly
```

```
## [1] "33.333%"
```

```
percdiff(3,4) # Dropping the arguments in directly, not named
```

```
## [1] "33.333%"
```

```
percdiff(second = 4, first = 3) # When named, order doesn't matter
```

```
## [1] "33.333%"
```

```
percdiff(second = 4,3) # Naming one, not the other
```

```
## [1] "33.333%"
```

```
percdiff(s = 4, f = 3) # Abbreviating the names (particularly dangerous)
```

```
## [1] "33.333%"
```

Writing functions

For readability and reliability I strongly suggest using the complete named form:

```
percdiff(first = 3, second = 4) # Naming the arguments explicitly  
  
## [1] "33.333%"
```

Writing functions

A function to find the most frequently occurring level of a factor:

```
mostoften <- function(x){  
  result <- sort(table(x),decreasing=TRUE)[1]  
  return(result)  
}
```

Fake some data to illustrate use of the function

```
set.seed(20160328)
```

#The "sample" command below will be the workhorse of Monte Carlo

```
a <- as.factor(sample(x=LETTERS,size=1000,replace=TRUE))
```

```
mostoften(x=a) # the argument named "x", takes on the value "a".
```

```
## T
```

```
## 51
```

Default arguments

A function can have default argument values assigned in the argument list. Note the additional argument, `n`, with the default value 1:

```
mostoften <- function(x,n=1){ #Pick off the top n items, default = 1 item
  result <- sort(table(x),decreasing=TRUE)[1:n]
  return(result)
}
mostoften(x=a) #Called with the default n = 1

##  T
## 51

mostoften(x=a,n=3) #Called with the n = 3

##  x
##  T  W  B
## 51 49 46
```

Checking the arguments

- It usually is a good idea to look at the values being passed into a function, to see if they make sense for the function to evaluate.
- If they don't make sense, then the options are either to stop the function, or to give a warning message and then coerce the values to something that does make sense.
- It's your choice (and that of other programmers) what to do, so you need to be aware of what the function may automatically do to its arguments.

Stopping the function

I only want this function to evaluate if the argument is a factor.
Stopping the function:

```
mostoften <- function(x,n=1){ #Pick off the top n items, default = 1  
#The next line will look at the argument to see if it is a factor variable  
#We haven't discussed "if" yet, but will do so soon.  
if(!is.factor(x)){ #The "is.factor" function returns TRUE or FALSE  
  stop("RW: This function is only meant to work on factors!")  
}  
result <- sort(table(x),decreasing=TRUE)[1:n]  
return(result)  
}  
b <- rnorm(100) #Create a numeric variable and try out the function  
mostoften(x = b, n = 3)  
  
## Error in mostoften(x = b, n = 3): RW: This function is only meant to  
work on factors!
```

Creating a warning

Making a warning instead:

```
mostoften <- function(x,n=1){ #Pick off the top n items, default = 1
  if(!is.factor(x)){
    warning("RW: This function is only meant to work on factors!
           I am going to coerce it to one.")
    x <- as.factor(x) #Coercion
  }
  result <- sort(table(x),decreasing=TRUE)[1:n]
  return(result)
}
b <- rnorm(100)#Create a numeric variable and try out the function
mostoften(x = b, n = 3)

## Warning in mostoften(x = b, n = 3): RW: This function is only meant
to work on factors!
##           I am going to coerce it to one.

## x
## -2.51072755618174 -2.11789073001765 -2.09755016836162
##           1           1           1
```


Environment and scope

- **Environment:** a collection of objects like data frames, functions etc.
- R has a nested set of environments. The top-level environment is called the global environment. R-Studio will show the objects in each environment.
- When you create a new function, it has its own environment.
- **Scope:** where and when can objects be seen in a program?
- A function can find variables in its own environment, plus those environments in which it is embedded.
- The consequence is that a function can use variables that are found in the global environment, without them being passed directly into the function.
- I strongly suggest, you do not do this because it can lead to odd behaviour that is very hard to debug!
- Good practice: use arguments for all the variables that will be used in your function. You will have much more predictable results.

Environment and scope

The following type of function works in R (though is NOT recommended):

```
a <- 10
f1 <- function(x){ #This is only a function of x
  x + a
}
f1(x = 5) #This function finds "a" in the global environment

## [1] 15

a <- 20 #But, some joker changes a, and now
#the function no longer gives the same answer as before,
#even though the function call is identical.
f1(x = 5)

## [1] 25
```

The main reason **we** want to understand this, is that it can be the root cause of a lot of bugs.

Returning many objects from a function

- It is quite possible that you want to return a variety of objects from a function, not just a single one.
- The natural way to do this is to create a list inside the function and then to return the whole list.
- We'll create a list that contains the most frequent level of a factor, and the rarest one.

Returning many objects from a single function

```
mostleast <- function(x){  
  #The first element in the sorted table  
  least.freq <- sort(table(x))[1]  
  #The last element in the sorted table  
  most.freq <- sort(table(x))[length(table(x))]  
  #Now make the list and return it  
  return(list(LEAST = least.freq, MOST = most.freq))  
}  
  
set.seed(20160328) #Create some data again  
a <- as.factor(sample(x=LETTERS,size=1000,replace=TRUE))  
mostleast(x = a)  
  
## $LEAST  
## 0  
## 27  
##  
## $MOST  
## T  
## 51
```

The three dots (ellipses), ..., special argument

Three dots. From the Wikipedia

Ellipsis (plural ellipses; from the Ancient Greek: "omission" or "falling short") is a series of dots (typically three, such as "...") that usually indicates an intentional omission of a word, sentence, or whole section from a text without altering its original meaning.

One useful feature of R is that we can have functions use other functions as arguments. We saw the `apply` function in HW 1:

```
my.matrix <- matrix(rnorm(100),ncol=10)
apply(X = my.matrix, MARGIN = 1, FUN = mean)

## [1] -0.33910110  0.02502231  0.17363863  0.20993464
## [5] -0.60397164 -0.26479405  0.10010886 -0.30265925
## [9] -0.38024362 -0.01646073
```

This applies the function `mean` to the rows of `my.matrix`.

The three dots (ellipses), ..., special argument

The help page for the `apply` function:

Usage

```
apply(X, MARGIN, FUN, ...)
```

Note the "...". These are the ellipses.

The three dots (ellipses), ..., special argument

We'll add in an NA to the matrix and repeat the apply function call:

```
my.matrix[1,1] <- NA # make the first element an NA
apply(X = my.matrix, MARGIN = 1, FUN = mean)

## [1] NA 0.02502231 0.17363863 0.20993464
## [5] -0.60397164 -0.26479405 0.10010886 -0.30265925
## [9] -0.38024362 -0.01646073
```

Notice that there is now an NA in the first element of the returned vector because the mean function fails on an NA.

However, there is an argument to the mean function: "na.rm" which determines how it should treat NAs. If we say "na.rm = TRUE" it will simply ignore the NAs and return the mean of the other values.

But how do we get the na.rm = TRUE into the mean function when using apply?

That's where the "..." comes in.

The three dots (ellipses), ..., special argument

Using the ellipses argument:

```
#Using the ... argument.
```

```
apply(X = my.matrix, MARGIN = 1, FUN = mean, na.rm = TRUE)
```

```
## [1] -0.25296617  0.02502231  0.17363863  0.20993464  
## [5] -0.60397164 -0.26479405  0.10010886 -0.30265925  
## [9] -0.38024362 -0.01646073
```

There is an additional argument to `mean` called `trim` which sets the fraction of data to be trimmed from either end before calculating the mean. This can be passed in through the ... too:

```
#Passing multiple arguments to the "mean" function
```

```
apply(X = my.matrix, MARGIN = 1, FUN = mean, na.rm = TRUE, trim = 0.1)
```

```
## [1] -0.25296617  0.04620525  0.22868084  0.23078294  
## [5] -0.61912342 -0.21812883  0.04875700 -0.15690288  
## [9] -0.40481106  0.03649638
```


Topics covered today include:

- Writing functions.
- Argument checking.
- Stop and warning.
- Environment and scope.
- The ellipsis argument.

Next time

- More on writing your own functions (examples) (II)

Today's function list

Do you know what each of these functions does?

```
... (ellipses)
body
environment
formals
function
glm
if
is.factor
lm
mean
predict
read.csv
return
sample
sprintf
stop
summary
warning
```